
Fixed-Priority Multiprocessor Scheduling: To Partition or not to Partition

Stefan Voigt

- ☐ Real-time scheduling vs. minimal makespan
 - Scheduling algorithms
 - Schedulability tests
- ☐ Periodic tasks vs. aperiodic tasks
- ☐ Fixed priority vs. dynamic priority
- ☐ Preemptive
- ☐ Partitioned vs. non-partitioned (global)
 - Comparison of both methods
 - Hybrid algorithms

-
- ☐ Introduction
 - Characteristics of both methods (partitioned + global)
 - ☐ Comparison of both methods
 - Introduction of some representatives
 - Performance
 - Preemption density
 - ☐ Conclusion

- ☐ Decision whether a task set is schedulable is NP-hard
- ☐ No method dominates the other
- ☐ But:

Partitioned method

- Provide performance guarantees
- Good average case performance
- Polynomial time (sufficient schedulability test)

☐ Non-partitioned method

- Received much less attention
- No efficient schedulability test exist (pessimistic)
- No efficient priority-assignment scheme has been found

- ☐ Two parts:
 - Dividing the task set into m groups
 - Scheduling each group locally on one processor
- ☐ The problem of scheduling each group of tasks on a processor is known
 - Rate monotonic scheduling (static priorities)
 - Earliest deadline first (dynamic priorities)
- ☐ Dividing the tasks into groups is NP-hard

Partitioned algorithm: RM-FFDU

Rate-Monotonic-First-Fit-Decreasing-Utilization

- Sort the task set (non-increasing utilization)
- Start with one processor
- For each task:
 - Try to assign the task τ_i to the processors P , starting with P_1
 - A task τ_i with utilization u_i can be assigned to P_j when:
$$u_i \leq 2 / \prod_{l=1}^{k_j} (u_{j,l} + 1) - 1 \quad (k_j = \text{number of tasks assigned to } P_j)$$
 - If τ_i cannot be assigned to the existing m processors, m will be increased by one and τ_i will be assigned to P_m

Partitioned algorithm: RM-FFDU

Example

$u_1=0.6$

$u_2=0.5$

$u_3=0.4$


$u_4=0.3$

$u_5=0.2$

$u_6=0.1$

$$u_i \leq 2 / \prod_{l=1}^{k_j} (u_{j,l} + 1) - 1$$


$$u \leq 2 / (1.6 // 1.5 // 1.4 // 1.3 // 1.2 // 1.1) - 1 = 0.02$$




Partitioned algorithm: RM-FFDU

Evaluation

- Worst-case tight bound of $5/3$ (comparison with optimal scheduling)
- Worst case:
 - $n = 15 \cdot k$ ($k \in \mathbb{N}$)
 - $u_i = 0.2 \ \forall i \in [1; n]$
 - RM-FFDU: 3 tasks per processor, $15k/3 = 5k$ processors



- Optimal scheduling: 5 tasks per processor, $15k/5 = 3k$ processors



- ☐ Worst-case tight bound of 2?
- ☐ Worst case:
 - $n = 2 \cdot k$ ($k \in \mathbb{N}$)
 - $u_i = 0.5 \quad \forall i \in [1; n]$
 - RM-FFDU: 1 tasks per processor, $2k/1 = 2k$ processors



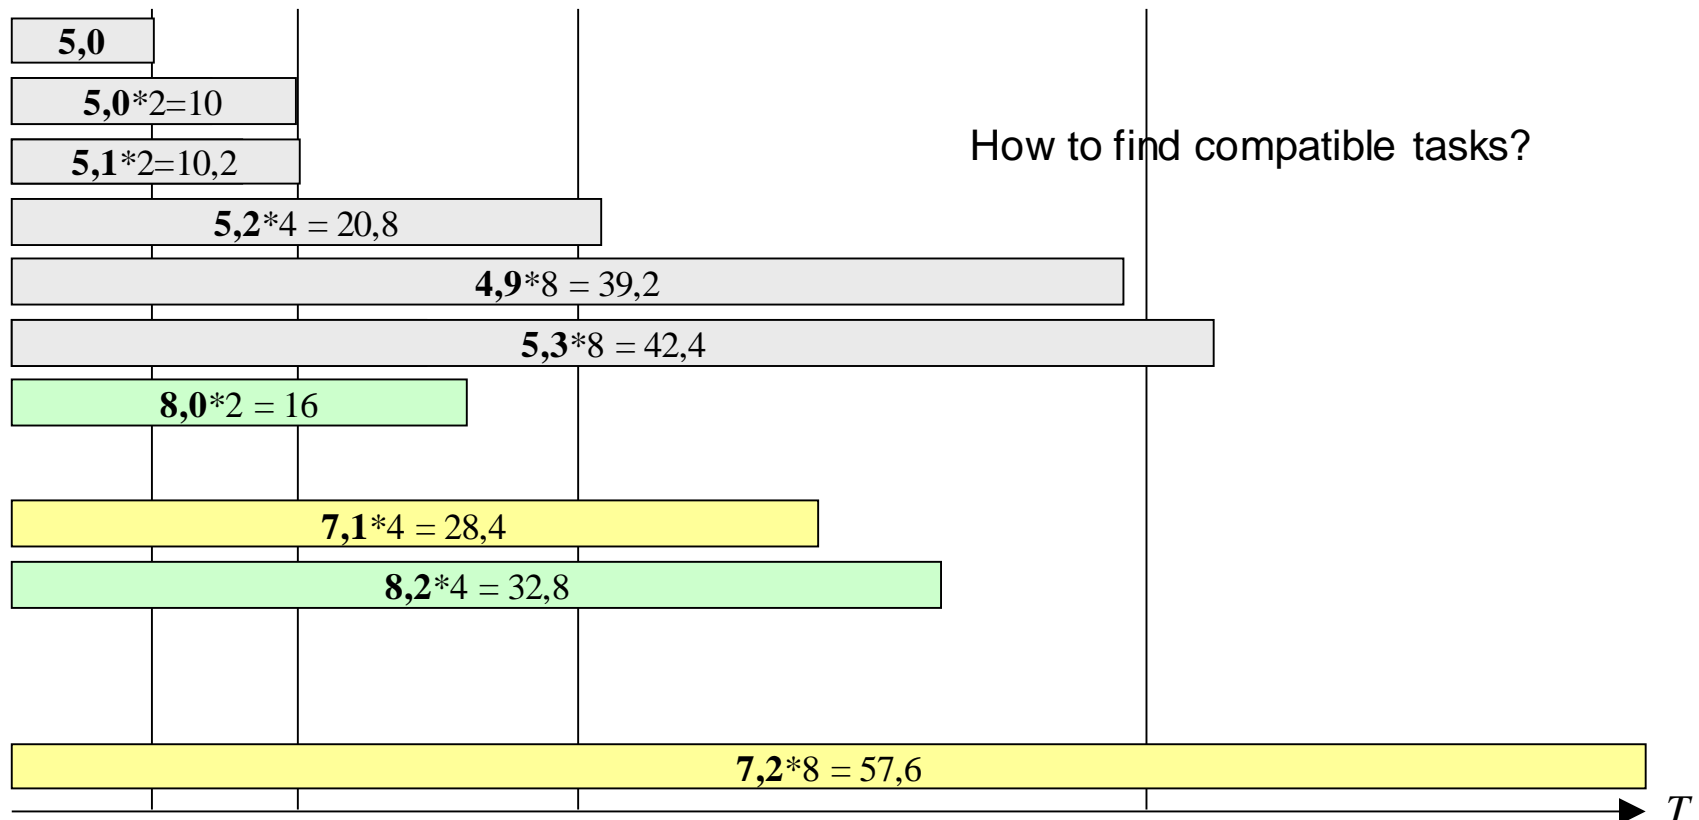
- Optimal scheduling: 2 tasks per processor, $2k/2 = k$ processors



- ☐ $O(nm + n \log n)$

Partitioned algorithm: R-BOUND-MP Compatible Tasks

- ☐ Tasks with same period
- ☐ Tasks that have a period closest to a power of two



Partitioned algorithm: R-BOUND-MP


Scale Task Set

□ $T = \{(C_1, T_1), (C_2, T_2), \dots, (C_n, T_n)\} \rightarrow T' = \{(2C_1, 2T_1), (C_2, T_2), \dots, (C_n, T_n)\}$

▪ u_i constant

▪ T' schedulable $\rightarrow T$ schedulable

Scaling factor: $s_i = 2^{\lfloor \lg_2 T_n / T_i \rfloor}$



-
- ☐ Main problem:
finding a priority assignment that guarantees schedulability as long as the system utilization is below a certain value

- Rate monotonic priority assignment
 - $p_i = 1/T_i$
- Suffers from Dhall's effect
 - Non-schedulable Task Set:
 - $T = \tau_1 \dots \tau_{m+1}$
 - $(T_i = 1, C_i = 2\varepsilon) \forall i \in [1; m+1]$
 - $(T_{m+1} = 1+\varepsilon, C_i = 1)$
 - τ_{m+1} has lowest priority and will miss its deadline
 - $\lim_{\varepsilon \rightarrow 0} U = 1$
 - $\lim_{m \rightarrow \infty, \varepsilon \rightarrow 0} U_S = 0$
- $O(n \log n)$ (sorting the task set)

- Guarantees that all task sets with $U_s \leq \text{US-LIMIT}$ are schedulable
- Tasks divided into two categories:
 - Tasks τ_i for which $U_i \leq \text{US-LIMIT}$:
 - $p_i = 1/(1+T_i)$ $p_i \in]0;1[$
 - Tasks τ_i for which $U_i > \text{US-LIMIT}$
 - $p_i = 1$
- Optimal US-LIMIT = 0.37482
- $O(n \log n)$ (sorting the task set)

-
- ☐ $p_i = 1 / (T_i - k * C_i)$
 - ☐ $U_s = 2 \frac{m}{3m - 1 + \sqrt{5m^2 - 6m + 1}}$
 - ☐ $\lim_{m \rightarrow \infty} U_s > 0.38$
 - ☐ $O(n \log n)$ (sorting the task set)

$$k = \frac{1}{2} \frac{m - 1 + \sqrt{5m^2 - 6m + 1}}{m}$$

- ☐ Partition as many tasks as possible on the given number of processors (RM-FFDU)
- ☐ Assign global priorities to the remaining tasks (adaptiveTkC)
- ☐ m local queues + 1 global queue
- ☐ If the local ready queue of a processor is empty, a task from the global ready queue is executed.

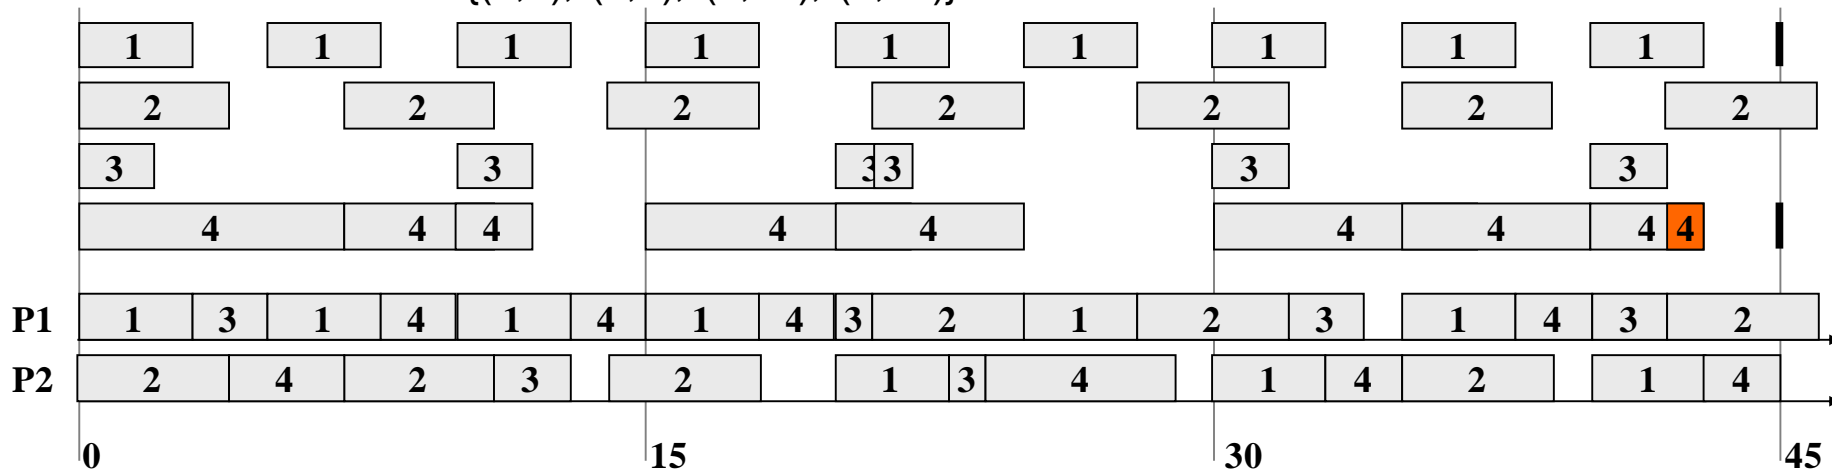
Performance comparison: Experimental setup

- ☐ $m = 4$ processors
- ☐ n : uniform distribution
 - $E[n] = 8$, minimum = $0.5 E[n]$, maximum = $1.5 E[n]$ ($n \in \{4, \dots, 12\}$)
- ☐ $T \in \{100, 200, \dots, 1500, 1600\}$
- ☐ u_i : normal distribution
 - $E[u_i] = 0.5$, $\text{stddev}[u_i] = 0.4$
 - $u_i < 0$ or $u_i > 1$: generate new u_i
- ☐ e_i : computed
 - $e_i = \text{floor}(u_i * t_i)$
 - $e = 0$: task generated again
- ☐ Success ratio
 - Fraction of all generated task sets that are successfully scheduled
 - For each point in a plot: average of 2,000,000 task sets

Performance comparison: Experimental setup


□ Task set is schedulable, when:

- Partitioned:
 - $m_{\text{required}} \leq m_{\text{given}}$
- Non-partitioned + hybrid:
 - Simulation of a meta period = $\text{LCM}(T_i) + \max(T_i)$
 - All task instances completed no later than their deadlines
- Why meta period?
 - $T = \{(3,5), (4,7), (2,10), (7,15)\}$




Performance comparison: Results


$[E[n]=8, E[u]=0.5, \text{stddev}[u]=0.4]$



Performance comparison: Results




Performance comparison: Results



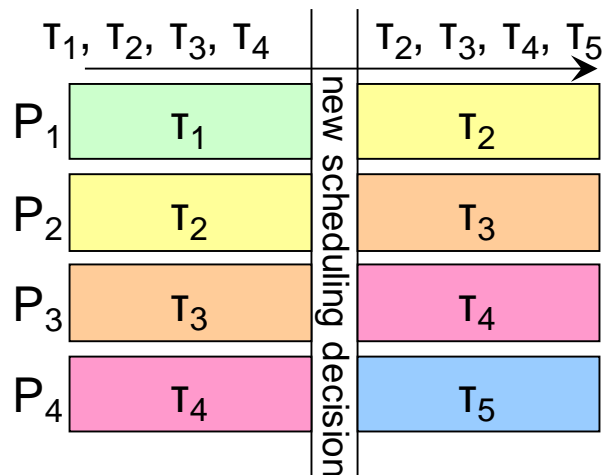
Performance comparison: Results

[m=4,E[n]=8,E[u]=0.5]

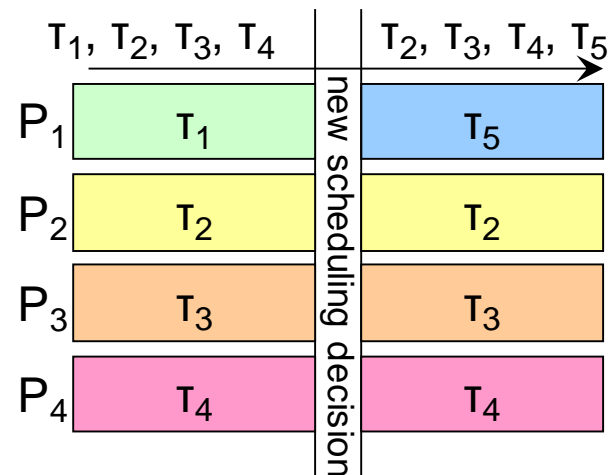


□ Architectural impact

- Costs for preemption and migration are not negligible
- Preemption aware Dispatcher for adaptiveTkC
- adaptiveTkCaware




Preemption unaware
Dispatcher

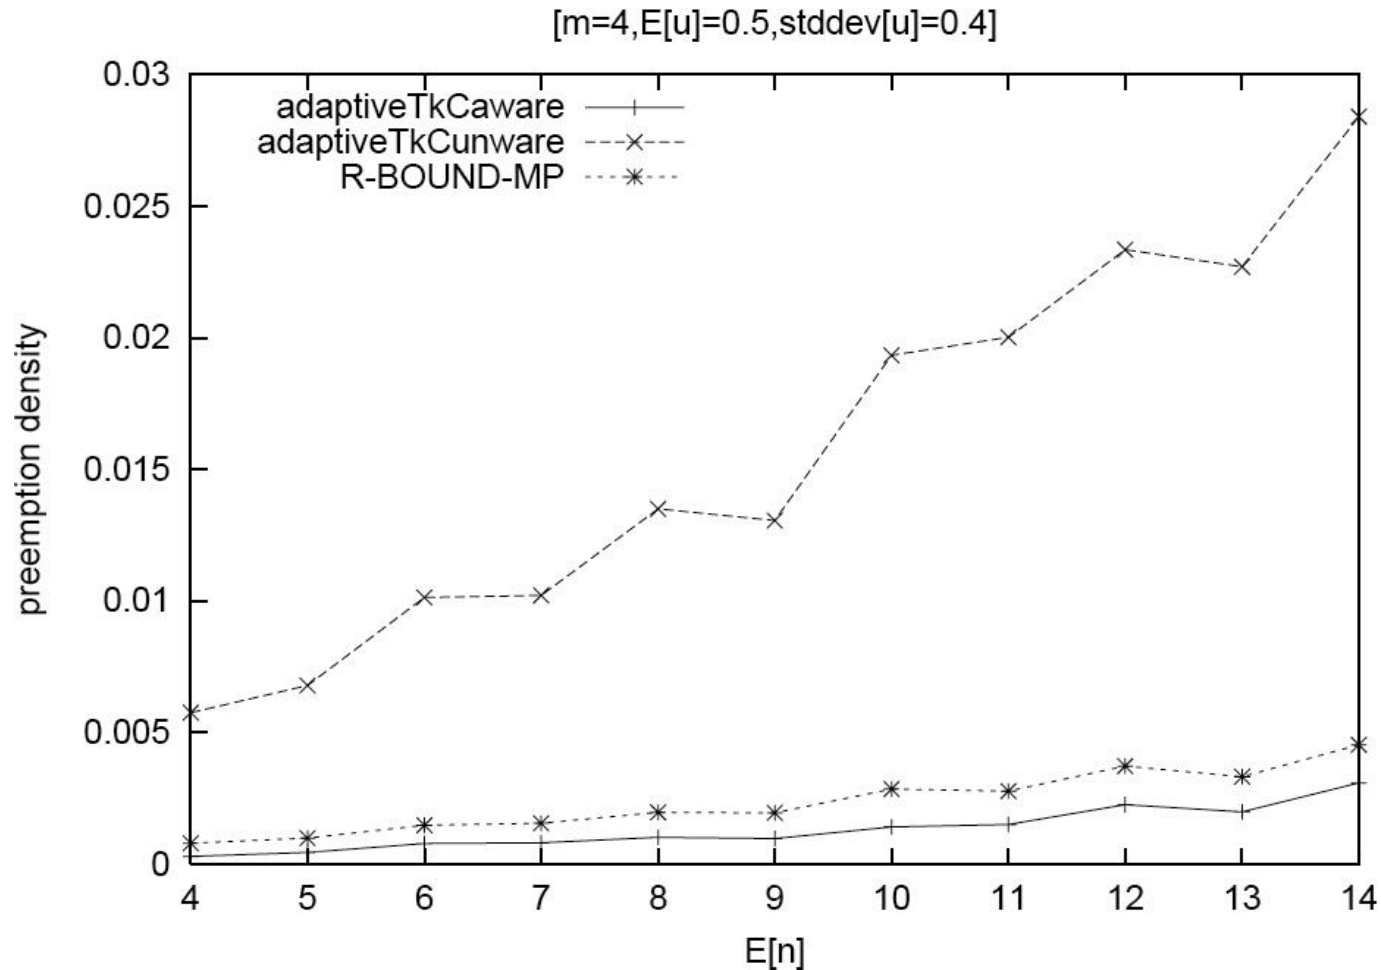


Preemption aware
Dispatcher


Preemption density comparison: Results



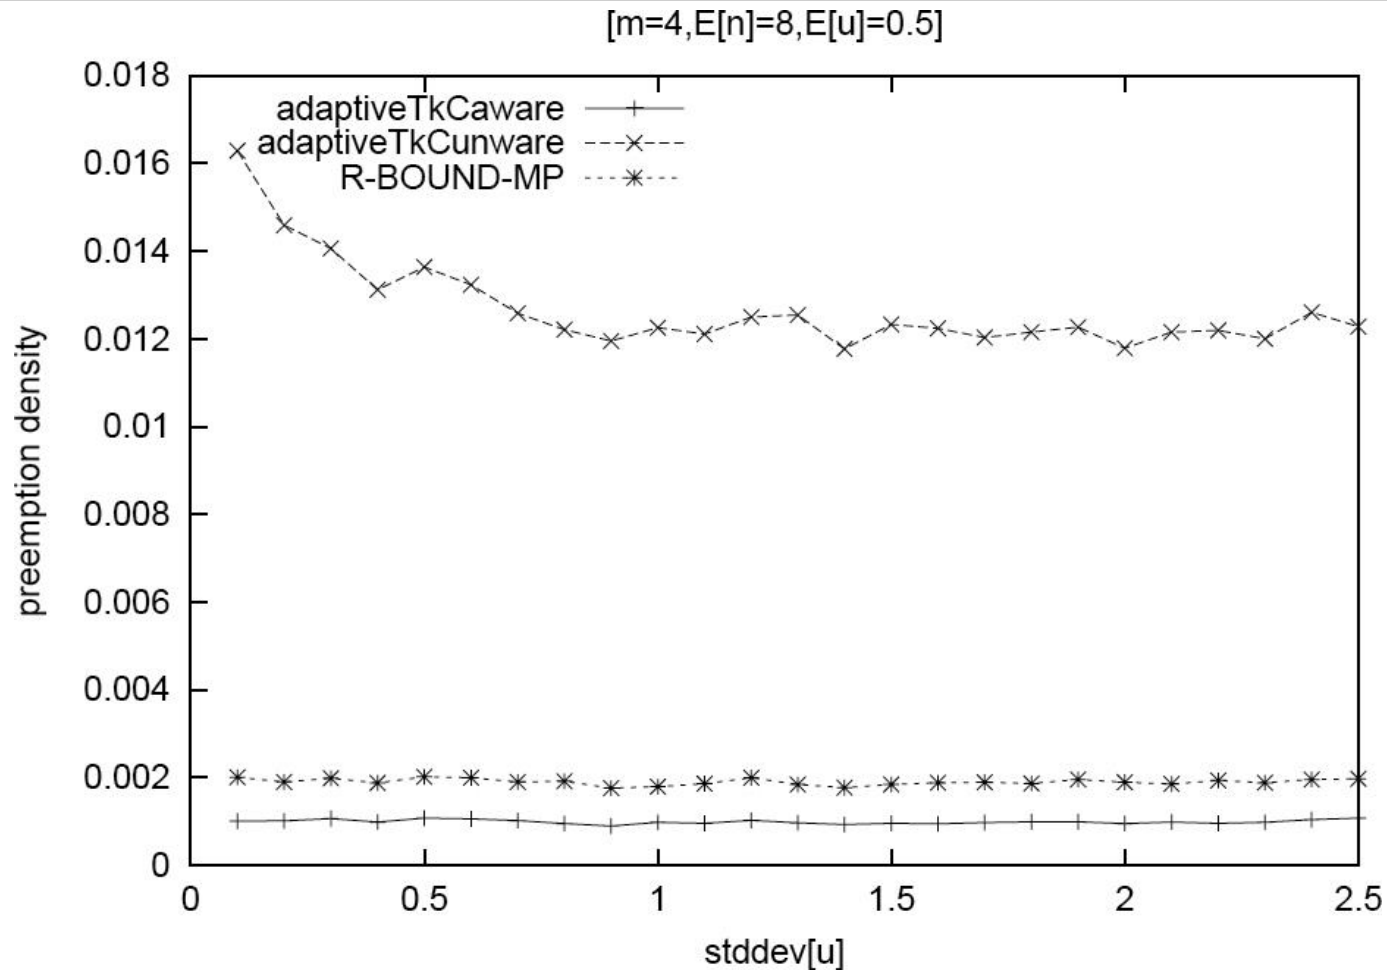
Preemption density comparison: Results



Preemption density comparison: Results



Preemption density comparison: Results



- System utilization:
 - Non-partitioned approach: $U_S < 0.38$
 - Partitioned approach: $U_S < 0.41$ (RM)
 - Varying execution times can cause low system utilization
- Computational complexity:
 - Non-partitioned approach: $O(n \log n)$
 - Partitioned approach: $O(nm + n \log n)$
- Preemption cost:
 - Non-partitioned approach can reduce preemptions using a preemption aware dispatcher

-
- [PartOrNot]** B. Andersson and J. Jonsson. **Fixed-priority preemptive multiprocessor scheduling: To partition or not to partition.** In Proc. of the International Conference on Real-Time Computing Systems and Applications, pages 337–346, Cheju Island, Korea, December 12–14, 2000.
- [RM]** C. L. Liu and J.W. Layland. **Scheduling algorithms for multiprogramming in a hard-real-time environment.** *Journal of the Association for Computing Machinery*, 20(1):46–61, January 1973.
- [RM-US[US-LIMIT]]** Lars Lundberg: **Analyzing Fixed-Priority Global Multiprocessor Scheduling.** IEEE Real Time Technology and Applications Symposium 2002: 145-153
- [AdTkC]** B. Andersson and J. Jonsson. **Some insights on fixed-priority preemptive non-partitioned multiprocessor scheduling.** In *Proc. of the IEEE Real-Time Systems Symposium – Workin- Progress Session*, Orlando, Florida, November 27–30, 2000. Also in TR-00-10, Dept. of Computer Engineering, Chalmers University of Technology.

-
- [RM-FFDU]** Y. Oh and S. H. Son. **Fixed-priority scheduling of periodic tasks on multiprocessor systems.** Technical Report 95-16, Department of Computer Science, University of Virginia, March 1995.
- [R-Bound-MP]** S. Lauzac, R. Melhem, and D. Moss´e. **An efficient RMS admission control and its application to multiprocessor scheduling.** In Proc. of the IEEE Int’l Parallel Processing Symposium, pages 511–518, Orlando, Florida, March 1998.
- [U bounds]** D. Oh and T. P. Baker. **Utilization bounds for n-processor rate monotone scheduling with static processor assignment.** Real-Time Systems, 15(2):183–192, September 1998.
- [Comp]** S. Lauzac, R. Melhem, and D. Moss´e. **Comparison of global and partitioning schemes for scheduling rate monotonic tasks on a multiprocessor.** In *10th Euromicro Workshop on Real Time Systems*, pages 188–195, Berlin, Germany, June 17–19, 1998.